Agent Policy Coverage and Attack Path

Which policy clauses address which risks, how an attack gets through, and what each defense buys

William Caban · 2026-10-24 · v1.0.0

demo · policy-coverage · attack-flow · synthetic-data

Agent: `support-agent (synthetic data)` · **Run:** `coverage-run-2026-10-24` · **Generated:** 2026-10-24T04:10:35Z

> Every value in this report is synthetic. It shows what three chart types add to a report: how a policy maps to risks, how an attack moves, and how several measurements compare at once.

1. Policy to risk mapping

Each curve joins a policy clause to a risk it addresses. Clauses and risks sit on one circle, grouped by policy section and risk category, and curves between groups are pulled together so shared routes read as bands. **A risk with no curve ending on it is not covered by any clause.** Here that is prompt injection.

Figure 1. Policy clauses and the risks they address

18 nodes, 15 links.

2. Attack path

A prompt-injection attack that ends in credential exfiltration, with the controls that interrupt it. Arrows point from an attack step to the next step, or to the risk it realizes; a control points at the step it blocks. The layout is computed from the links, so position carries no meaning: follow the arrows.

Figure 2. Attack flow and the controls that interrupt it

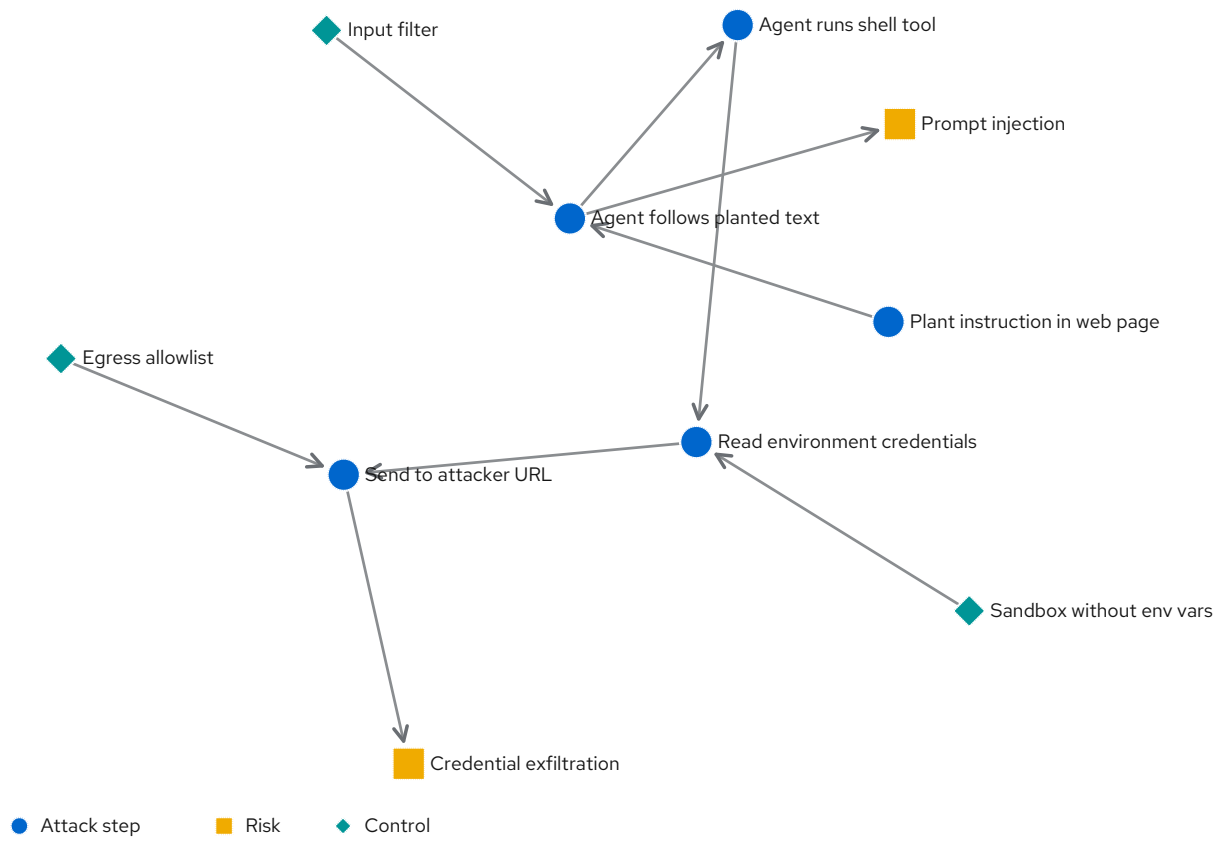

10 nodes, 9 links. Arrows point to the target. Layout is computed from the links; position carries no meaning.

3. What each defense buys

Each line is one agent configuration and each vertical axis is one measurement on its own scale. Axes where lower is better are flipped, so a line near the top is good on every axis. The sandbox run has no latency measurement: its line stops at the previous axis and is not drawn through zero.

Figure 3. Four agent configurations across five measurements

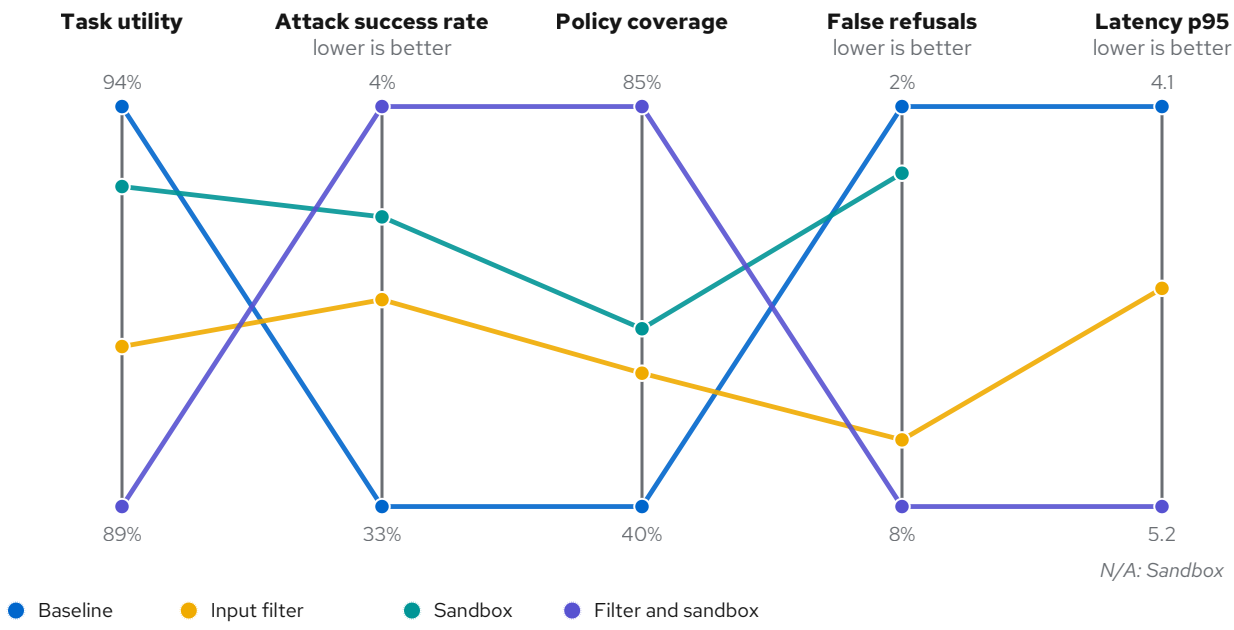

Each axis has its own scale, shown at both ends. Axes marked lower is better are flipped: up is always better.
N/A: 1 value not available; the line breaks there and is not drawn through zero.

The exact values behind the lines:

Configuration	Task utility	Attack success	Policy coverage	False refusals	p95 latency (s)
Baseline	94%	33%	40%	2%	4.1
Input filter	91%	18%	55%	7%	4.6
Sandbox	93%	12%	60%	3%	*data unavailable*
Filter and sandbox	89%	4%	85%	8%	5.2

Evidence

Spec `policy-coverage` v1.0.0 (sha256:ad047a384ec8), generated 2026-10-24T04:10:35Z. Verdicts come from the producers of the results, not from this renderer.

Fonts: Red Hat Text and Red Hat Mono (SIL OFL 1.1). Status icons: Font Awesome Free by Fonticons, Inc. (CC BY 4.0).

Source	Adapter	SHA-256	Evaluations
policy-coverage.jsonl	jsonl	`55df783cf4858d8f`	coverage, attack, configs